


Mean recorded billed cost across paid models


● LLM ■ *CodeAct ◆ RLM

Recorded cost per generated trajectory (CHF)


(a) Tier 1: Structural lookup ($n=10$)




(b) Tier 2: Property aggregation ($n=20$)



(c) Tier 3: Reaction-level reasoning ($n=35$)



(d) Tier 4: Relational reasoning ($n=35$)



Paid models only: Gemini, GPT-5 mini, and Claude. Free SwissAI models are excluded; failed-attempt cost metadata are unavailable.