


RLM resource use grows with task and corpus scale



RLM only; unweighted mean $\pm$ SEM across Qwen, Gemini, GPT-5 mini, and Claude. Additive metrics summarize 5,930 successful scored trajectories from 1,791 jobs; peak memory is measured per successful job. Recorded latency and process wall time are independent elapsed-time measurements.