


Qwen 3.5





LLM CodeAct RLM



Terminal failed jobs score zero; unresolved jobs are excluded from the current curve.

DeepSeek V4 Flash


LLM *CodeAct *RLM



Terminal failed jobs score zero; unresolved jobs are excluded from the current curve. * Arm has running, stale, pending, or unrun jobs in the gold snapshot.

GLM 5.2





● LLM ■ *CodeAct ◆ *RLM



Terminal failed jobs score zero; unresolved jobs are excluded from the current curve. * Arm has running, stale, pending, or unrun jobs in the gold snapshot.

Gemini 3.7 Flash


LLM CodeAct RLM



Terminal failed jobs score zero; unresolved jobs are excluded from the current curve.

GPT-5 mini


LLM CodeAct RLM




Claude Haiku 4.5

LLM CodeAct RLM


(a) Tier 1: Structural lookup ($n=10$)




(b) Tier 2: Property aggregation ($n=20$)



(c) Tier 3: Reaction-level reasoning ($n=35$)



(d) Tier 4: Relational reasoning ($n=35$)



Terminal failed jobs score zero; unresolved jobs are excluded from the current curve.