



Mean wall time for successful answers

● LLM ■ CodeAct ◆ *RLM


(a) Tier 1: Structural lookup ($n=10$)




(b) Tier 2: Property aggregation ($n=20$)



(c) Tier 3: Reaction-level reasoning ($n=35$)



(d) Tier 4: Relational reasoning ($n=35$)



Unweighted mean $\pm$ SEM across terminal model arms. Only successful jobs enter the wall-time numerator and denominator; failed jobs are excluded.