


Mean answer time across models


● LLM ■ *CodeAct ◆ *RLM

Wall time per answered trajectory (s)


(a) Tier 1: Structural lookup ($n=10$)




(b) Tier 2: Property aggregation ($n=20$)



(c) Tier 3: Reaction-level reasoning ($n=35$)



(d) Tier 4: Relational reasoning ($n=35$)



Unweighted mean $\pm$ SEM across terminal model arms. Wall time covers answered trajectories.